



南京大學

NANJING UNIVERSITY

应用层

殷亚凤

智能软件与工程学院

苏州校区南雍楼东区225

yafeng@nju.edu.cn , <https://yafengnju.github.io/>



Internet Applications

- Internet Applications Overview
- WWW and HTTP
- Electronic Mail
- Domain Name Service (DNS)
- File Transfer Protocol (FTP)
- Content Distribution Networks (CDNs)





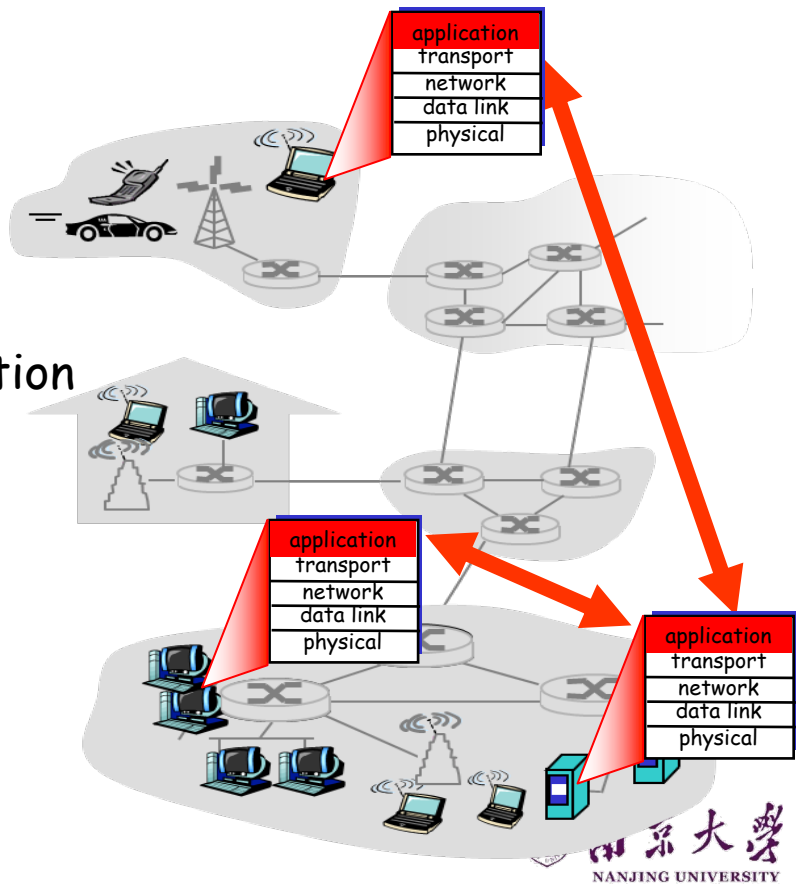
Internet Applications Overview

Application: communicating, distributed processes

- e.g., Email, Web, P2P file sharing, instant messaging
- Running in end systems (hosts)
- Exchange messages to implement application

Application-layer protocols

- One "piece" (agent) of an app
- Define messages exchanged by apps and actions taken
- Use communication services provided by lower layer protocols (TCP, UDP, RTP)





Application Architectures

possible structure of applications:

- client-server (CS)
- peer-to-peer (P2P)






Client-Server Paradigm

Client:

- Start as required
- Initiates contact with server, "**speaks first**"
- Host may have dynamic IP addresses
- e.g. Web: client implemented in browser;
Email: in mail reader

Server:

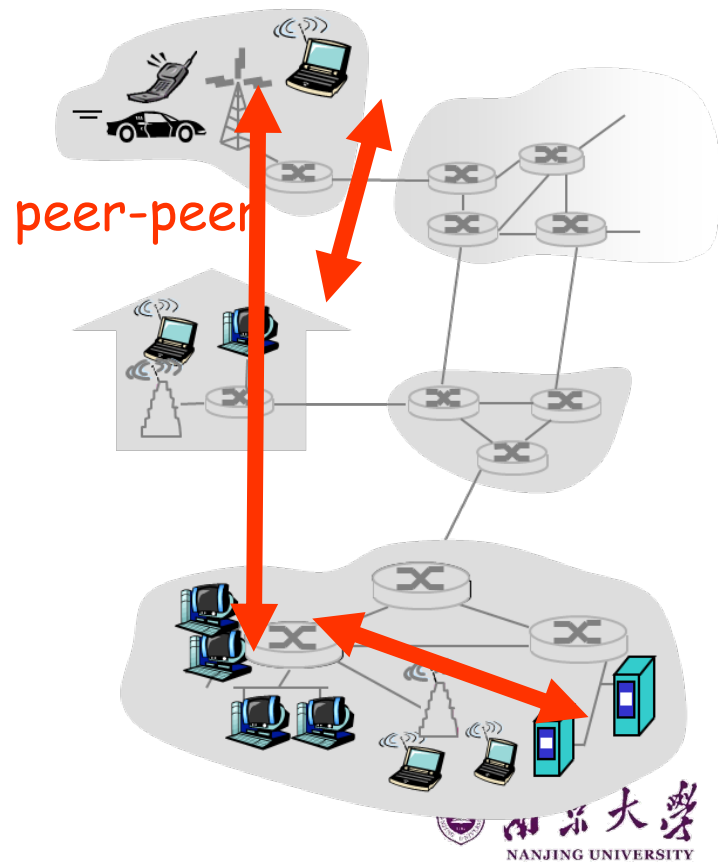
- Run as daemon (always-on)
- **Provides requested service** to Client
- Host has permanent IP address
- e.g. Web server sends requested Web page,
mail server delivers Email





Peer-to-Peer Paradigm

- No always-on server
- Arbitrary end systems directly communicate
- peers request service from other peers, provide service in return to other peers
 - self scalability - new peers bring new service capacity, as well as new service demands
- Peers are intermittently connected and change IP addresses
 - Highly scalable but difficult to manage
- Examples: Gnutella, BitTorrent, Skype





Client-Server and P2P

Skype

- Voice-over-IP P2P application
- **Centralized server:** finding address of remote party
- Direct client-client connection

Instant messaging

- Chatting between two users is P2P
- **Centralized service:** user presence detection/location
- User registers its IP address with central server when it comes online
- User contacts central server to find IP addresses of parties





Jargons of Internet Applications

- **Process**: program running within a host
 - Within same host, 2 processes communicate using **inter-process communication** (defined by OS)
 - Processes running in different hosts communicate with an **app-layer protocol**
- **User agent**: interfaces with app “above” and network “below”
 - Implements user interface & **app-layer protocol**, e.g.
 - Web: browser, web server
 - Email: mail reader, mail server
 - Streaming audio/video: media player, media server





Typical Applications

- Web and HTTP
- Email
- DNS
- FTP
- CDN





Internet Applications

- Internet Applications Overview
- WWW and HTTP
- Electronic Mail
- Domain Name Service (DNS)
- File Transfer Protocol (FTP)
- Content Distribution Networks (CDNs)





Web components

- Infrastructure:
 - Clients
 - Servers (DNS, CDN, Datacenters)
- Content:
 - URL: naming content
 - HTML: formatting content
- Protocol for exchanging information: HTTP





URL - Uniform Resource Locator

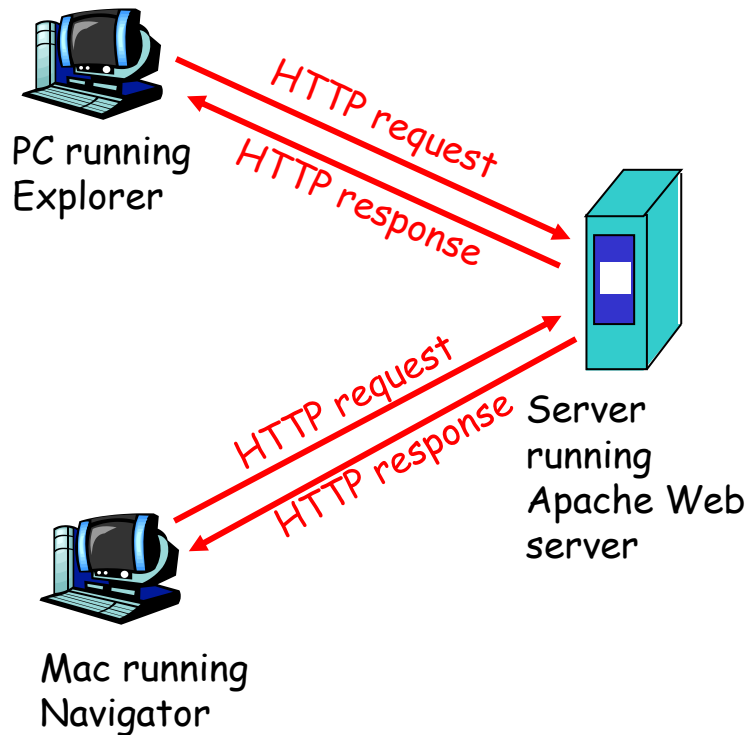
- A unique identifier for an object on WWW
- URL format
`<protocol>://<host>:<port>/<path>?query_string`
 - Protocol: method for transmission or interpretation of the object, e.g. http, ftp, Gopher
 - Host: DNS name or IP address of the host where object resides
 - Path: pathname of the file that contains the object
 - **Query_string**: name/value pairs sent to app on the server
- An example
`http://www.nju.edu.cn:8080/somedir/page.htm`





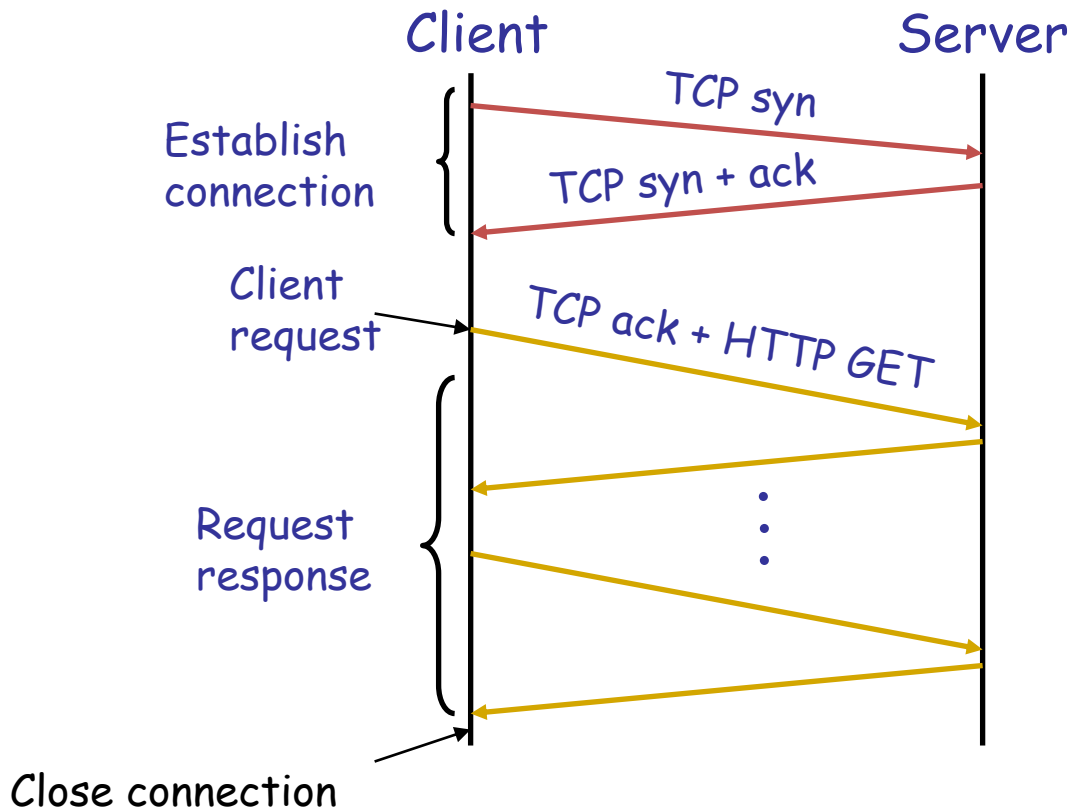
Hyper Text Transfer Protocol (HTTP)

- Client-server architecture
 - Server is "always on" and "well known"
 - Clients initiate contact to server
- Synchronous request/reply protocol
 - Runs over TCP, Port 80
- Stateless
- ASCII format
 - Before HTTP/2





Steps in HTTP request/response





Method types (HTTP 1.1)

- GET, HEAD
- POST
 - Send information (e.g., web forms)
- PUT
 - Uploads file in entity body to path specified in URL field
- DELETE
 - Deletes file specified in the URL field





Client-to-server communication

- HTTP Request Message
 - Request line: method, resource, and protocol version

The diagram illustrates the structure of an HTTP request. It shows a sequence of lines: a request line, followed by header lines, and a carriage return line feed. The request line is circled in blue and contains the text "GET /somedir/page.html HTTP/1.1". The header lines are grouped by a bracket and contain the text "Host: www.someschool.edu", "User-agent: Mozilla/4.0", "Connection: close", and "Accept-language: fr". A red arrow points from the text "carriage return line feed indicates end of message" to the end of the header lines.

```
request line → GET /somedir/page.html HTTP/1.1
header lines → Host: www.someschool.edu
               User-agent: Mozilla/4.0
               Connection: close
               Accept-language: fr
               (blank line)
carriage return line feed →
```

indicates end of message





Server-to-client communication

- HTTP Response Message

- Status line: protocol version, status code, status phrase
- Response headers: provide information
- Body: optional data

status line

(protocol, status code, status phrase)

header lines

data

e.g., requested HTML file

HTTP/1.1 200 OK

Connection close

Date: Thu, 06 Jan 2017 12:00:15 GMT

Server: Apache/1.3.0 (Unix)

Last-Modified: Mon, 22 Jun 2006 ...

Content-Length: 6821

Content-Type: text/html

(blank line)

data data data data data ...





HTTP is stateless

- Each request-response treated independently
 - Servers not required to retain state
- **Good:** Improves scalability on the server-side
 - Failure handling is easier
 - Can handle higher rate of requests
 - Order of requests doesn't matter
- **Bad:** Some applications need persistent state
 - Need to uniquely identify user or store temporary info
 - e.g., Shopping cart, user profiles, usage tracking, ...





Question

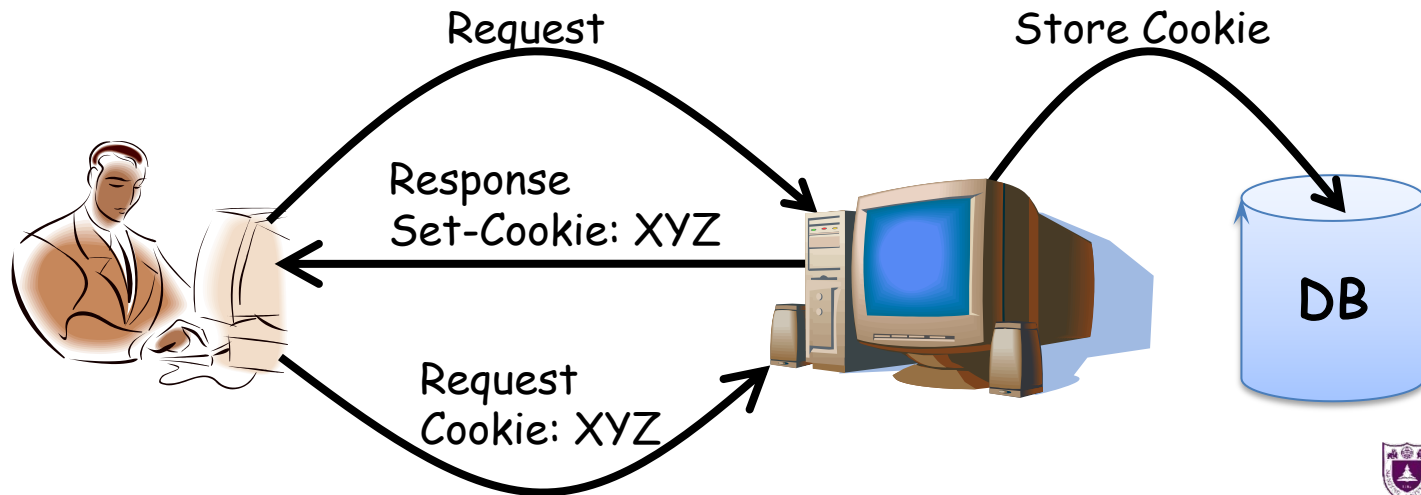
- How does a stateless protocol keep state?





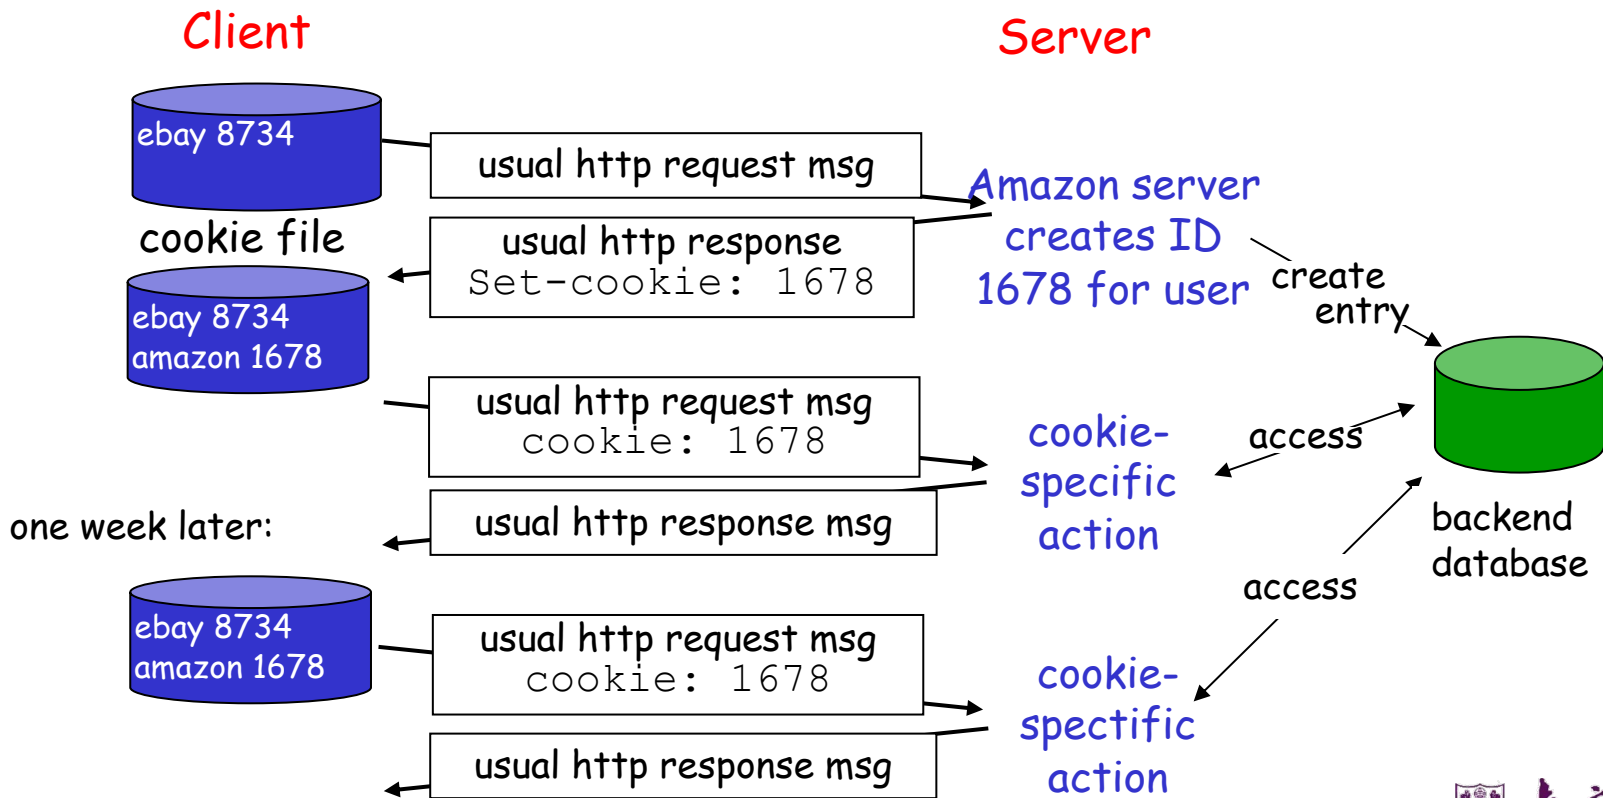
State in a stateless protocol: Cookies

- Client-side state maintenance
 - Client stores small state on behalf of server
 - Client sends state in future requests to the server
- Can provide authentication





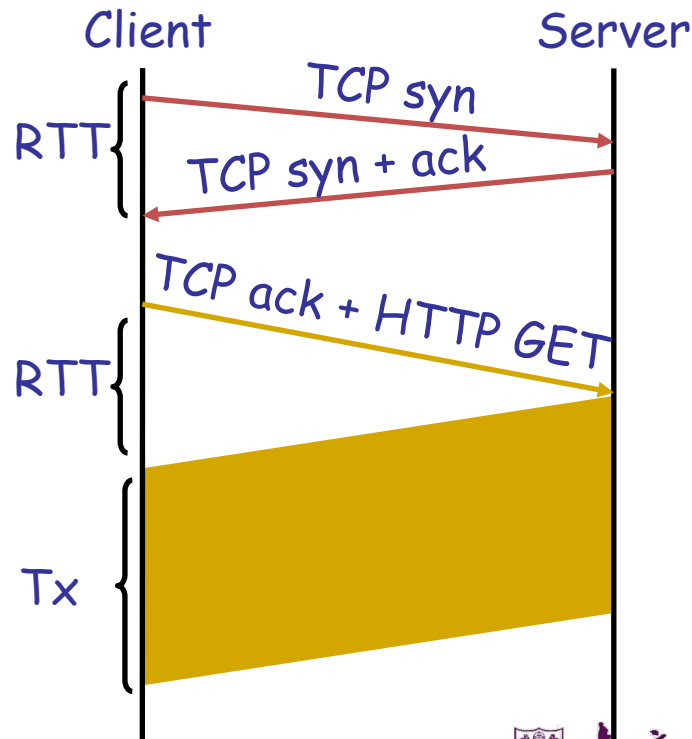
A Cookies Example





HTTP performance: Object request response time

- RTT (round-trip time)
 - Time for a small packet to travel from client to server and back
- Response time
 - 1 RTT for TCP setup
 - 1 RTT for HTTP request and first few bytes
 - Transmission time
 - **Total** = 2RTT + Transmission Time





Non-persistent connections

- Default in HTTP/1.0
- $2RTT + \Delta$ for each object in the HTML file!
 - One more $2RTT + \Delta$ for the HTML file itself
- Doing the same thing over and over again
 - Inefficient

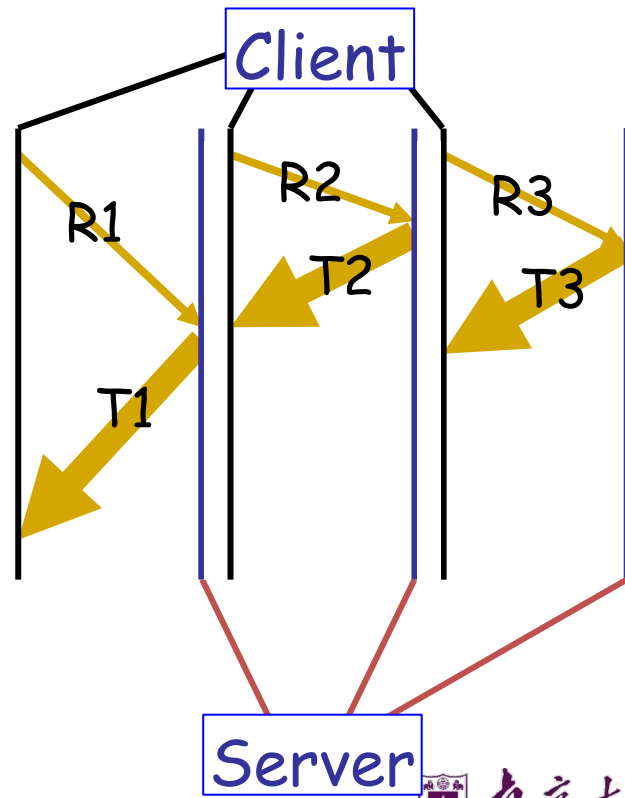




Concurrent requests and responses

- Use multiple connections in parallel
- Does not necessarily maintain order of responses

- Client = 😊
- Content provider = 😊
- Network = 😞 Why?





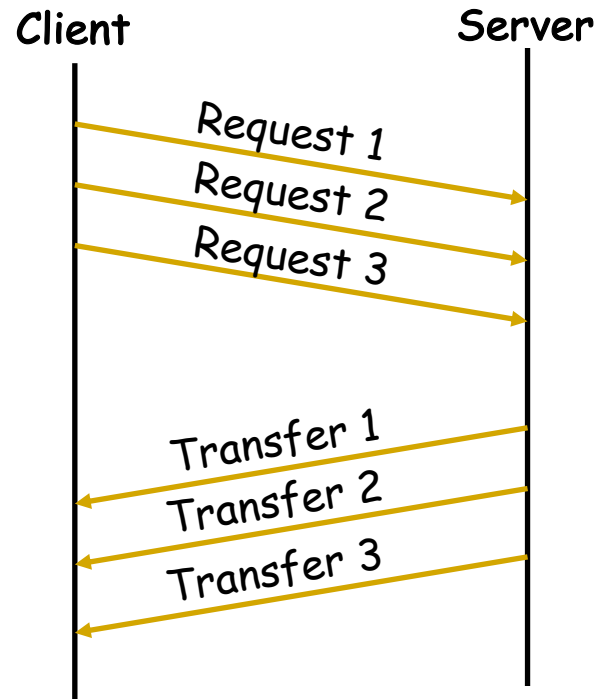
Persistent connections

- Maintain TCP connection across multiple requests
 - Including transfers subsequent to current page
 - Client or server can tear down connection
- Advantages
 - Avoid overhead of connection set-up and tear-down
 - Allow underlying layers (e.g., TCP) to learn about RTT and bandwidth characteristics
- Default in HTTP/1.1



Pipelined requests & responses

- Batch requests and responses to reduce the number of packets
- Multiple requests can be contained in one TCP segment





Caching

- Why does caching work?
 - Exploits locality of reference
- How well does caching work?
 - Very well, up to a limit
 - Large overlap in content
 - But many unique requests
 - ✓ A universal story!
 - ✓ Effectiveness of caching grows logarithmically with size





Caching: How

- Modifier to **GET** requests:
 - **If-modified-since** - returns “not modified” if resource not modified since specified time

GET /somedir/page.html HTTP/1.1

Host: www.someschool.edu

User-agent: Mozilla/4.0

If-modified-since: Wed, 18 Jan 2017 10:25:50 GMT
(blank line)





Caching: How

- Modifier to **GET** requests:
 - **If-modified-since** - returns “not modified” if resource not modified since specified time
- Client specifies “**if-modified-since**” time in request
- **Server** compares this against “**last modified**” time of resource
- **Server** returns “**Not Modified**” if resource has not changed
- or a “OK” with the latest version otherwise



Caching: How

- Modifier to GET requests:
 - If-modified-since - returns “not modified” if resource not modified since specified time
- Response header:
 - Expires - how long it's safe to cache the resource
 - No-cache - ignore all caches; always get resource directly from server



Caching: Where?


- Options
 - Client (browser)
 - Forward proxies
 - Reverse proxies
 - Content Distribution Network





Caching: Where?

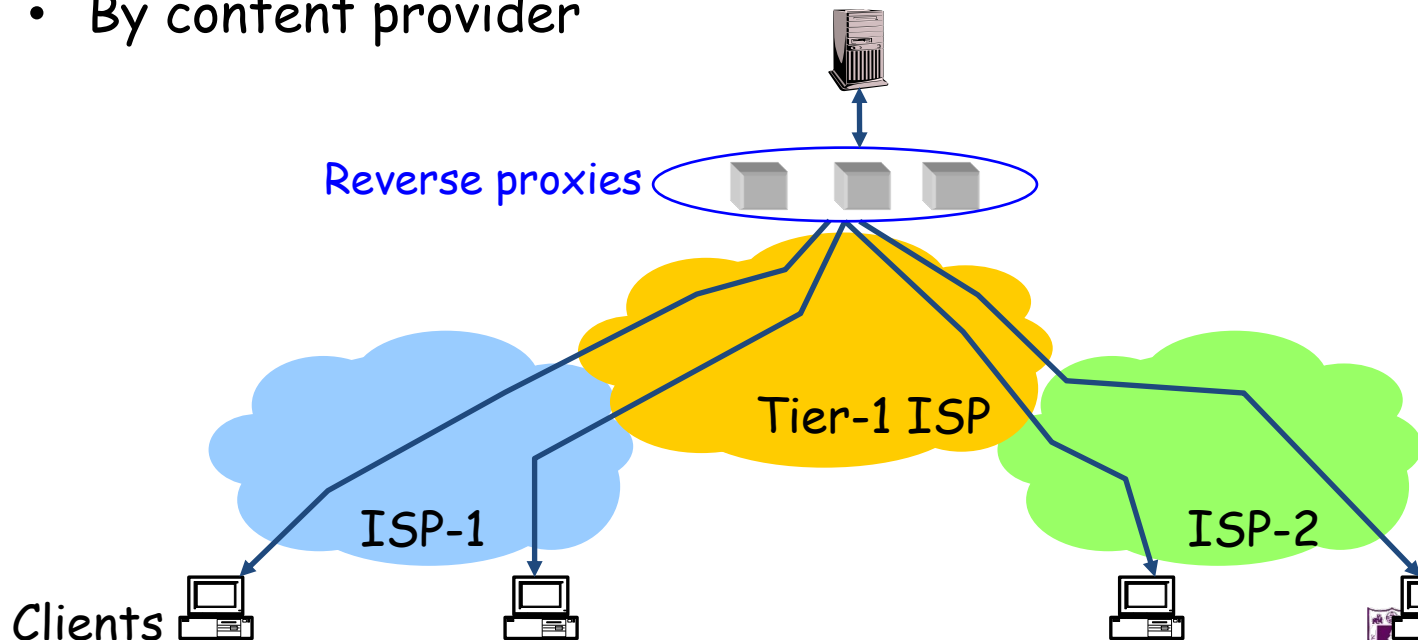
- Many clients transfer same information
 - Generate unnecessary server and network load
 - Clients experience unnecessary latency





Caching with Reverse Proxies


- Cache documents close to server
 - Decrease server load
 - By content provider





Caching with Forward Proxies

- Cache documents close to clients
 - Reduce network traffic and decrease latency
 - By ISPs or enterprises





Internet Applications

- Internet Applications Overview
- Domain Name Service (DNS)
- Electronic Mail
- File Transfer Protocol (FTP)
- WWW and HTTP
- Content Distribution Networks (CDNs)





Electronic Mail

- One of most heavily used apps on Internet
- **SMTP**: Simple Mail Transfer Protocol
 - Delivery of simple **text** messages
- **MIME**: Multi-purpose Internet Mail Extension
 - Delivery of other types of data, e.g. **voice, images, video clips**
- **POP**: Post Office Protocol
 - **Msg retrieval from server**, including authorization and download
- **IMAP**: Internet Mail Access Protocol
 - Manipulation of stored **msgs on server**






Components of Email System

User Agent

- Composing, editing, reading mail messages
- e.g. Eudora, Outlook, Foxmail, Netscape Messenger
- Outgoing, incoming mail messages stored on server

Mail Servers (Host)

- **Mailbox** contains incoming mail messages for user
- **Message queue** of outgoing mail messages
- **SMTP protocol** between mail servers to send mail messages





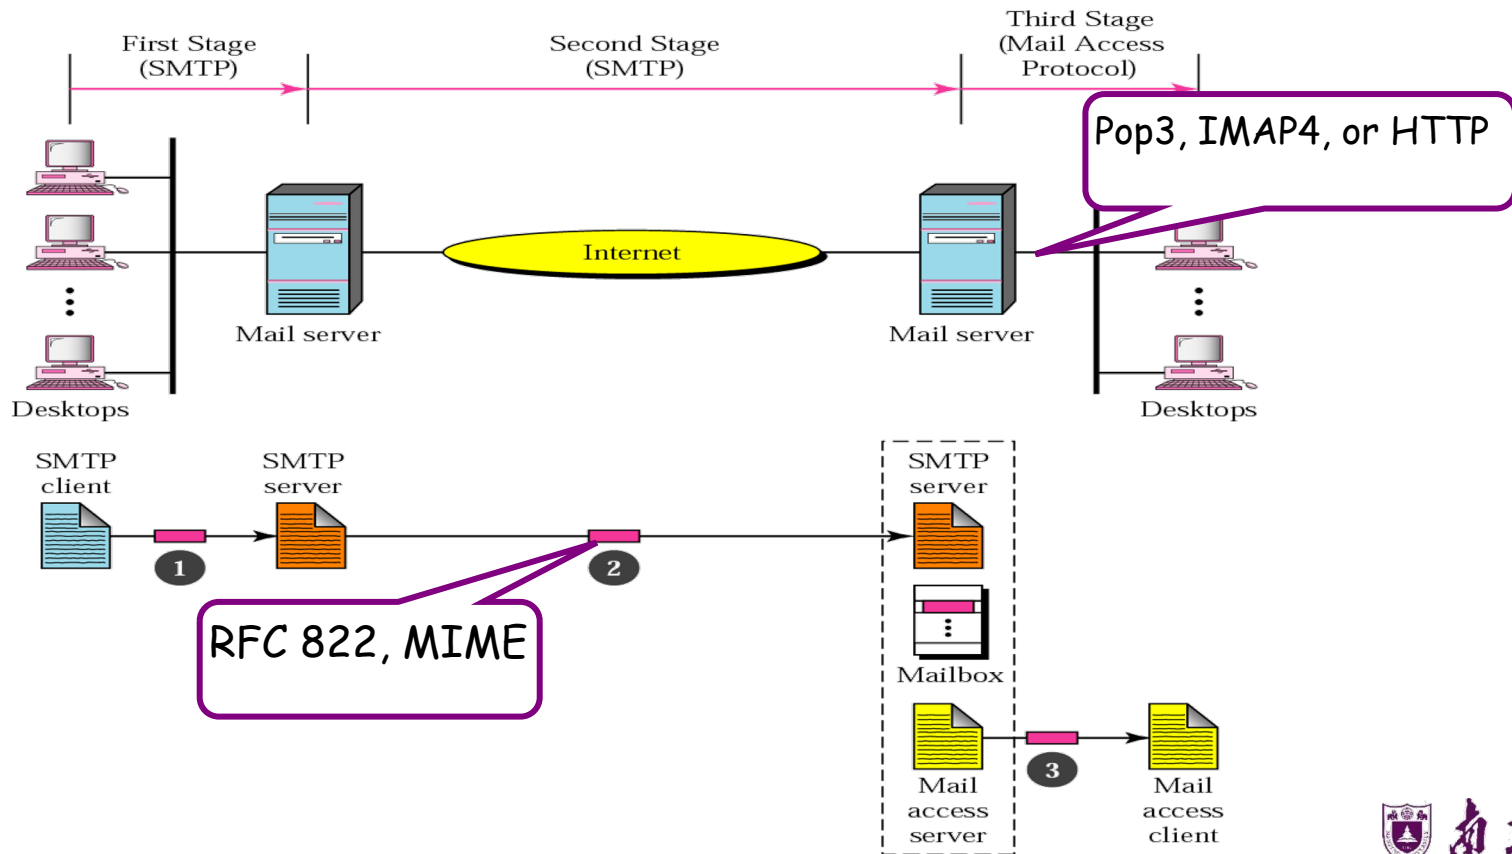
3 Stages of Mail Delivery

- 1st Stage
 - Email goes from **local user agent** to the **local SMTP server**
 - User agent acts as SMTP client
 - Local server acts as SMTP server
- 2nd Stage
 - Email is relayed by the local server to the **remote SMTP server**
 - Local server acts as SMTP client now
- 3rd Stage
 - The **remote user agent** uses a mail access protocol to access the mailbox on remote server
 - POP3 or IMAP4






Illustration of Mail Delivery





A Mail Delivery Scenario

- 1) Alice uses UA to compose a mail message and [to bob@some school.edu](mailto:bob@some school.edu)
- 2) Alice's UA sends mail to her mail server using SMTP, mail placed in [message queue](#)
- 3) Client side of SMTP opens TCP connection with Bob's mail server
- 4) SMTP client sends Alice's mail over the TCP connection
- 5) Bob's mail server places the mail in Bob's [mailbox](#)
- 6) Bob invokes his UA to read the mail, e.g. by Pop3





SMTP Transaction

3 phases of transfer

- Handshaking (greeting)
- Transfer of one or more mails data
- Close connection

Command/response interaction

- **Commands:** ASCII text
- **Response:** status code and phrase


```
S: 220 hamburger.edu
C: HELO crepes.fr
S: 250 Hello crepes.fr, pleased to meet you
C: MAIL FROM: <alice@crepes.fr>
S: 250 alice@crepes.fr ... Sender ok
C: RCPT TO: <bob@hamburger.edu>
S: 250 bob@hamburger.edu ... Recipient ok
C: RCPT TO: <Johm@hamburger.edu>
S: 550 No such user here
C: DATA
S: 354 Enter mail, end with "." on a line by itself
C: Do you like ketchup?
C:   How about pickles?
C: .
S: 250 Message accepted for delivery
C: QUIT
S: 221 hamburger.edu closing connection
```



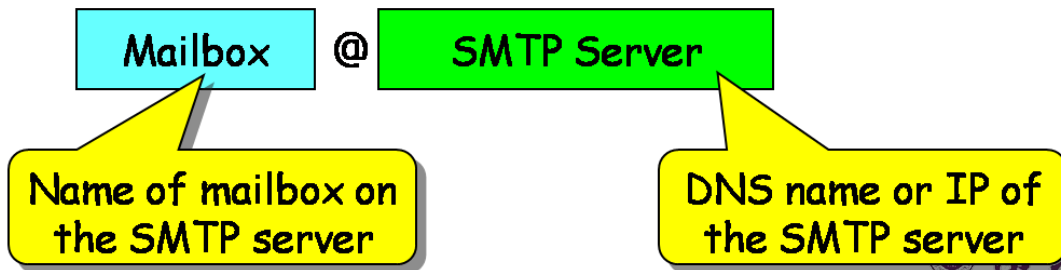


An Email Message

- Header lines, e.g.
 - To: Alice@sina.com
 - From: Bob@gmail.com
 - Subject: Dinner tonight
- Body
 - Mail contents, ASCII characters only



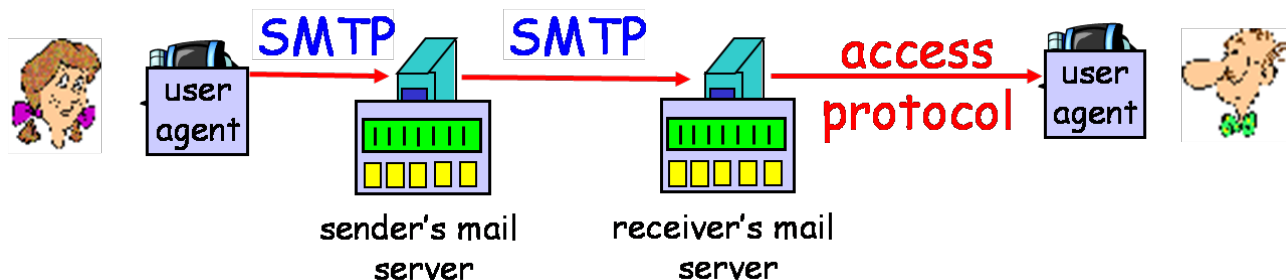
■ Mail destinations





Mail Access Protocols

- **SMTP**: delivery/storage to receiver's server
- **Mail access protocol**: mail retrieval from server
- **POP**: Post Office Protocol [RFC 1939]
 - Authorization (agent <-->server) and download
- **IMAP**: Internet Mail Access Protocol [RFC 1730]
 - more features, including manipulation of stored mails on server
- **HTTP**: gmail, Hotmail, Yahoo!, etc.





Internet Applications

- Internet Applications Overview
- WWW and HTTP
- Electronic Mail
- Domain Name Service (DNS)
- File Transfer Protocol (FTP)
- Content Distribution Networks (CDNs)





Domain Name Service (DNS)

- **Function**
 - Map “domain names” into IP addresses
 - e.g. www.baidu.com → 119.75.217.109
- **Domain Name System**
 - **Distributed database** implemented in hierarchy of many **name servers**
 - **App-layer protocol** host and name servers to communicate to resolve “domain names”
 - **Load balancing**: set of IP addresses for one server name

Q: why not centralize DNS?

- single point of failure
- traffic volume
- distant centralized database
- maintenance

A: doesn't scale!





Goals

- **Uniqueness**: no naming conflicts
- Scalable
 - Many names and frequent updates (secondary)
- Distributed, autonomous administration
 - Ability to update my own (machines') names
 - Don't have to track everybody's updates
- Highly available
- Lookups are fast
- Perfect consistency is a **non-goal**





How?

- Partition the namespace
- Distribute administration of each partition
 - Autonomy to update my own (machines') names
 - Don't have to track everybody's updates
- Distribute name resolution for each partition
- How should we partition things?



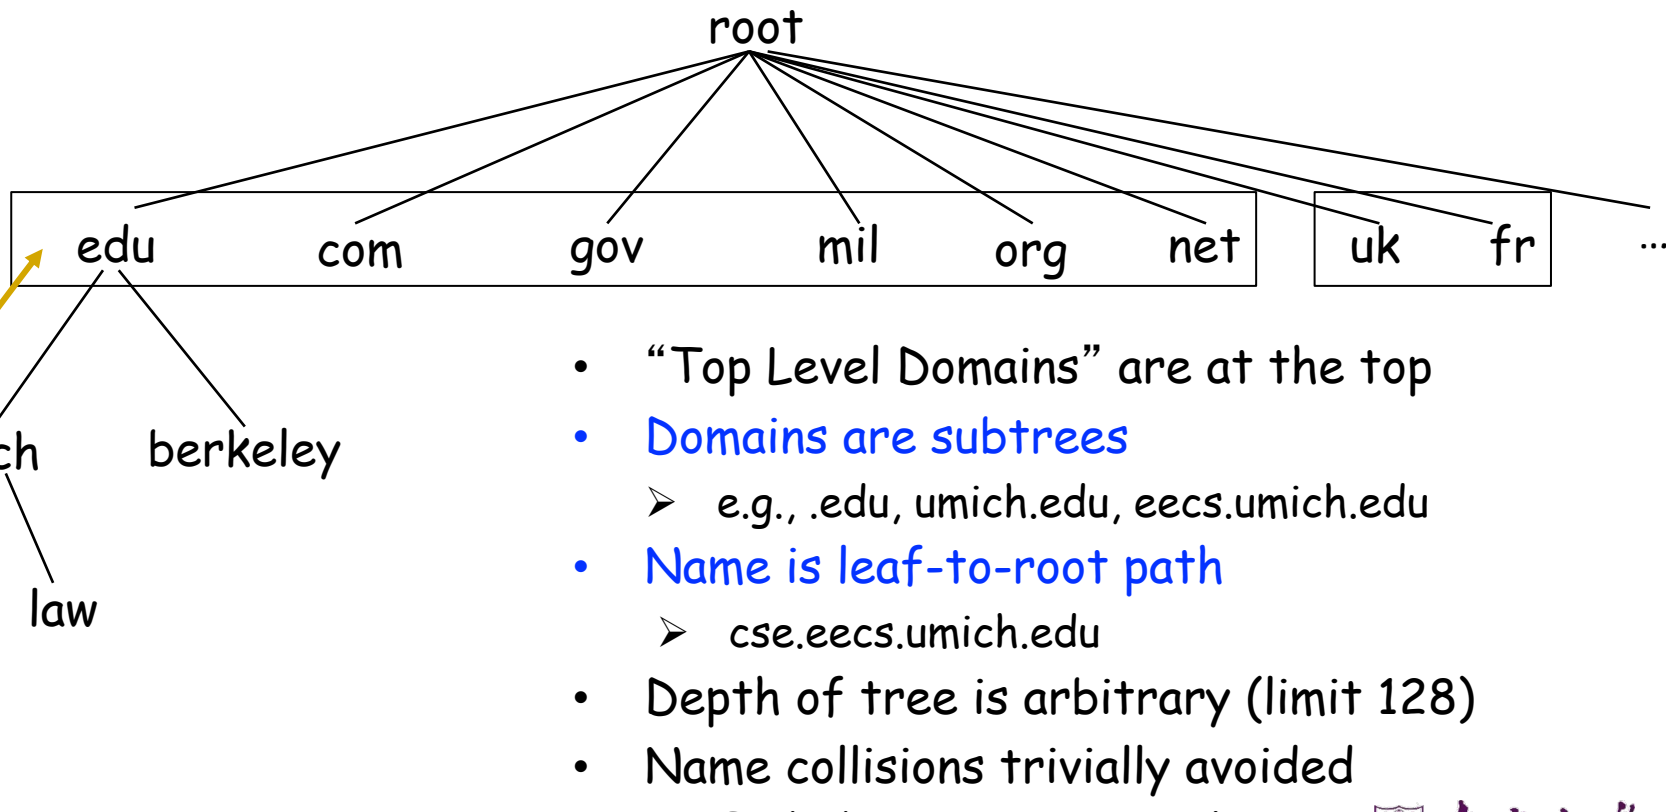
Key idea: Hierarchy

- Three intertwined hierarchies
 - Hierarchical namespace
 - ✓ As opposed to original flat namespace
 - Hierarchically administered
 - ✓ As opposed to centralized
 - (Distributed) hierarchy of servers
 - ✓ As opposed to centralized storage





Hierarchical namespace

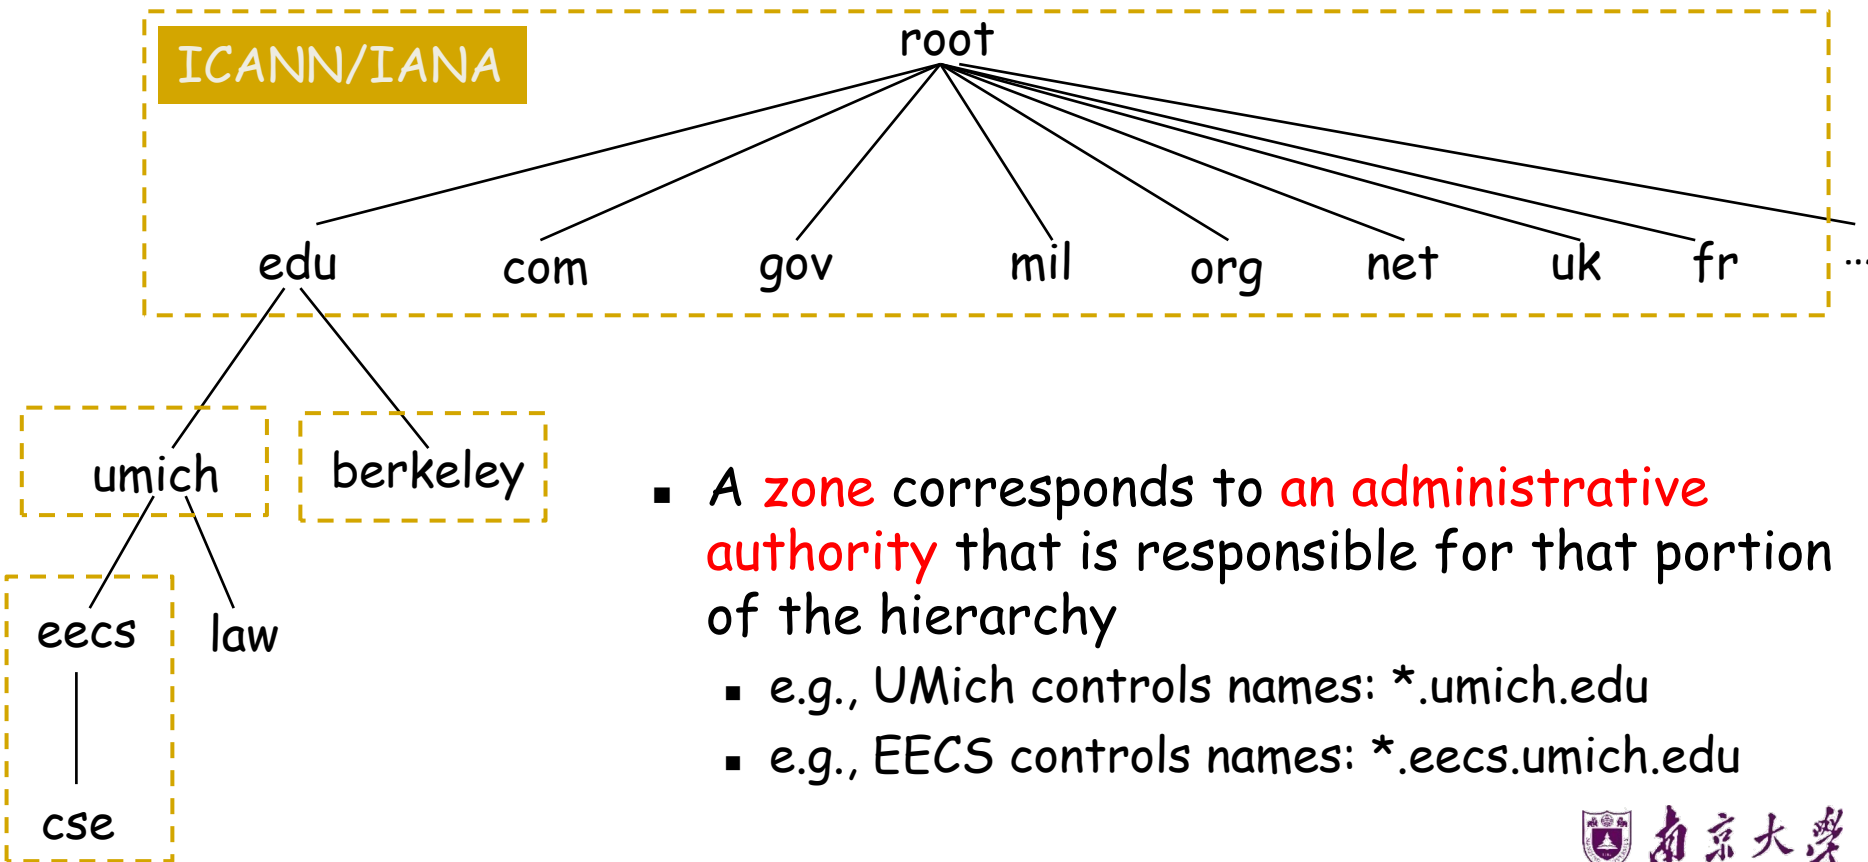


- “Top Level Domains” are at the top
- Domains are subtrees
 - e.g., .edu, umich.edu, eeecs.umich.edu
- Name is leaf-to-root path
 - cse.eeecs.umich.edu
- Depth of tree is arbitrary (limit 128)
- Name collisions trivially avoided
 - Each domain is responsible





Hierarchical administration



- A **zone** corresponds to an **administrative authority** that is responsible for that portion of the hierarchy
 - e.g., UMich controls names: *.umich.edu
 - e.g., EECS controls names: *.eeecs.umich.edu





Hierarchy of DNS Servers

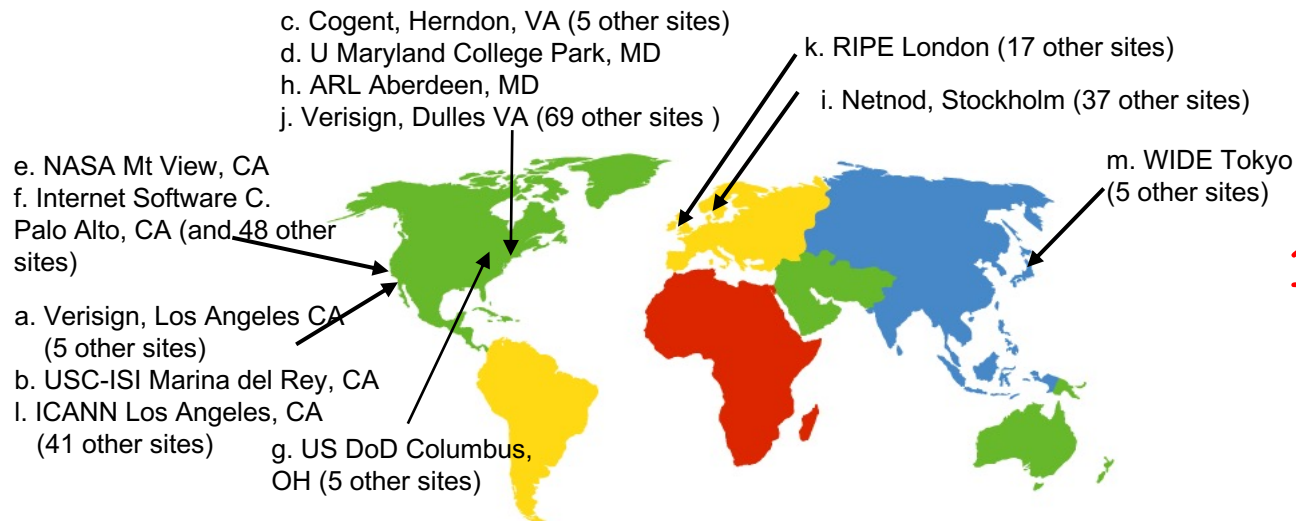
- **Root name servers**
 - Contacted by local name server that can not resolve name
- **Top-level domain servers**
 - Responsible for com, org, net, edu, etc, and all top-level country domains, e.g. cn, uk, fr
- **Authoritative DNS servers**
 - Organization's DNS servers, providing authoritative **hostname to IP mappings**
- **Local Name Servers**
 - Maintained by each residential ISP, company, university
 - When host makes DNS query, query is sent to its local DNS server





DNS: root name servers

- root name server:
 - returns IP mappings of TLD servers



13 root name
“servers”
worldwide





TLD, authoritative servers

- Top-level domain (TLD) servers:
 - responsible for com, org, net, edu, aero, jobs, museums, and all top-level country domains, e.g.: uk, fr, ca, jp
 - Network Solutions maintains servers for .com TLD
 - Educause for .edu TLD
- Authoritative DNS servers:
 - organization's own DNS server(s), providing authoritative **hostname to IP mappings** for organization's named hosts
 - can be maintained by organization or service provider






Local DNS name server

- Does not strictly belong to hierarchy
- Each ISP (residential ISP, company, university) has one
 - also called “default name server”
- When host makes DNS query, query is sent to its local DNS server
 - has local cache of recent name-to-address translation pairs (but may be out of date!)
 - acts as proxy, forwards query into hierarchy



DNS Name Resolution Example

- Bob at cis.poly.edu wants IP address for Alice at gaia.cs.umass.edu
- Iterated query:
- Contacted server replies with name of next server to contact
- Host-Server: recursive query
- Server-Server: iterative query





DNS Records

- A DNS resource record (RR)

RR format: (name, value, type, ttl)

- "Name" is the domain name, "type" denotes how "value" is explained
 - e.g. Name Server records (NS), Mail Exchangers (MX), Host IP Address (A), Canonical name (CNAME)
- Examples
 - (networkutopia.com, dns1.networkutopia.com, NS, 32768)
 - (dns1.networkutopia.com, 212.212.212.1, A, 5600)





DNS protocol

- Query and Reply messages; both with the same message format
 - Header: identifier, flags, etc.
 - Plus resource records
 - See text/section for details
- Client-server interaction on UDP Port 53
 - Spec supports TCP too, but not always implemented





DNS caching

- Performing all these queries takes time
 - Up to 1-second latency before starting download
- **Caching can greatly reduce overhead**
 - The top-level servers very rarely change
 - Popular sites (e.g., www.cnn.com) visited often
 - Local DNS server often has the information cached
- How DNS caching works
 - DNS servers **cache responses to queries**
 - Responses include a “**time to live**” (TTL) field
 - Server **deletes cached entry** after TTL expires





Internet Applications


- Internet Applications Overview
- Domain Name Service (DNS)
- Electronic Mail
- File Transfer Protocol (FTP)
- WWW and HTTP
- Content Distribution Networks (CDNs)





File Transfer Protocol (FTP)

- RFC 959, use TCP, port 21/20
- Transfer file to/from remote host
- Client/Server model, client side initiates file transfer (either to/from remote)
- Deals with **heterogeneous** OS and file systems
- Needs **access control** on remote file system





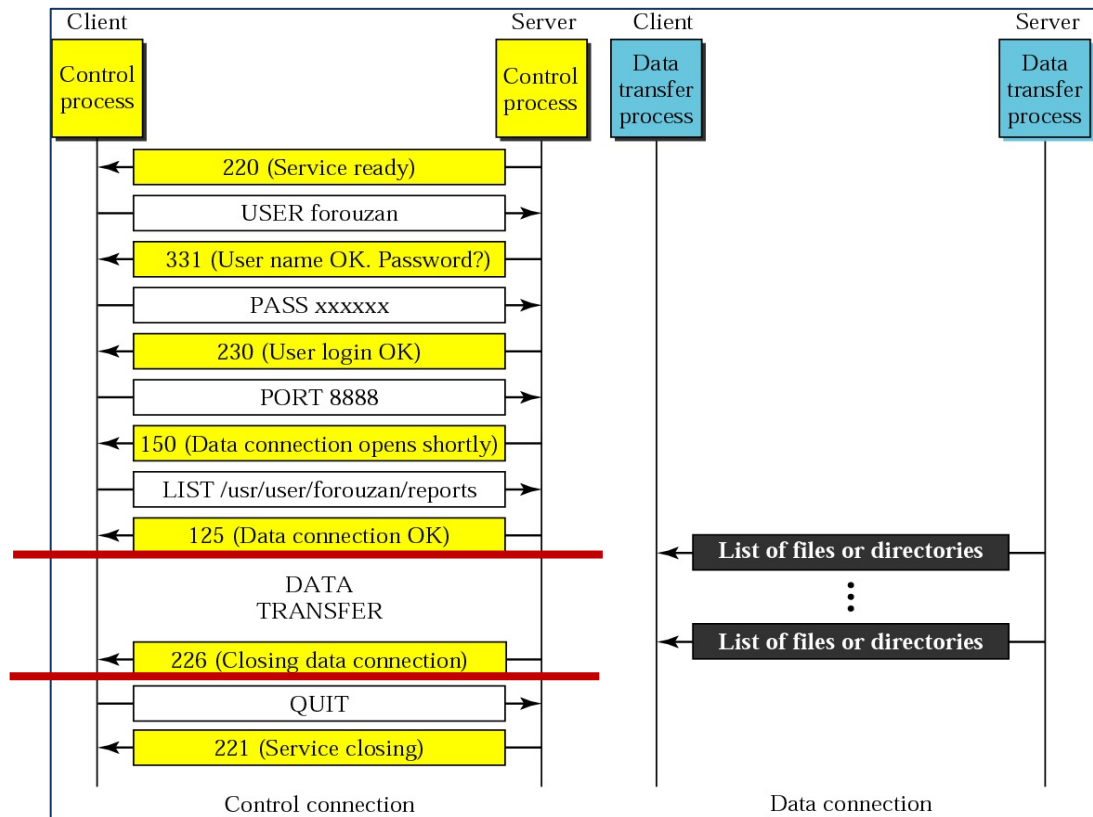
Control and Data Connections

- FTP client contacts FTP server at port 21, opens a **control connection**
- Client authorized over control connection
- Client browses **remote directory** by sending commands over control connection
- When server receives file transfer command, **server opens 2nd TCP data connection (for file) to client**
 - One connection for each file transferred
- After transferring one file, server closes data connection
- Control connection stays **"out of band"**
- FTP server maintains **"user state"**: current directory, earlier authentication





Illustration of FTP Session





FTP Commands and Responses

Sample commands:

- Sent as ASCII text over control channel
- USER username
- PASS password
- LIST return list of file in current directory
- RETR filename retrieves (gets) file
- STOR filename stores (puts) file onto remote server

Sample return codes:

- Status code and phrase (as in HTTP)
- 331 Username OK, password required
- 125 data connection already open; transfer starting
- 425 Can't open data connection
- 452 Error writing file





BitTorrent

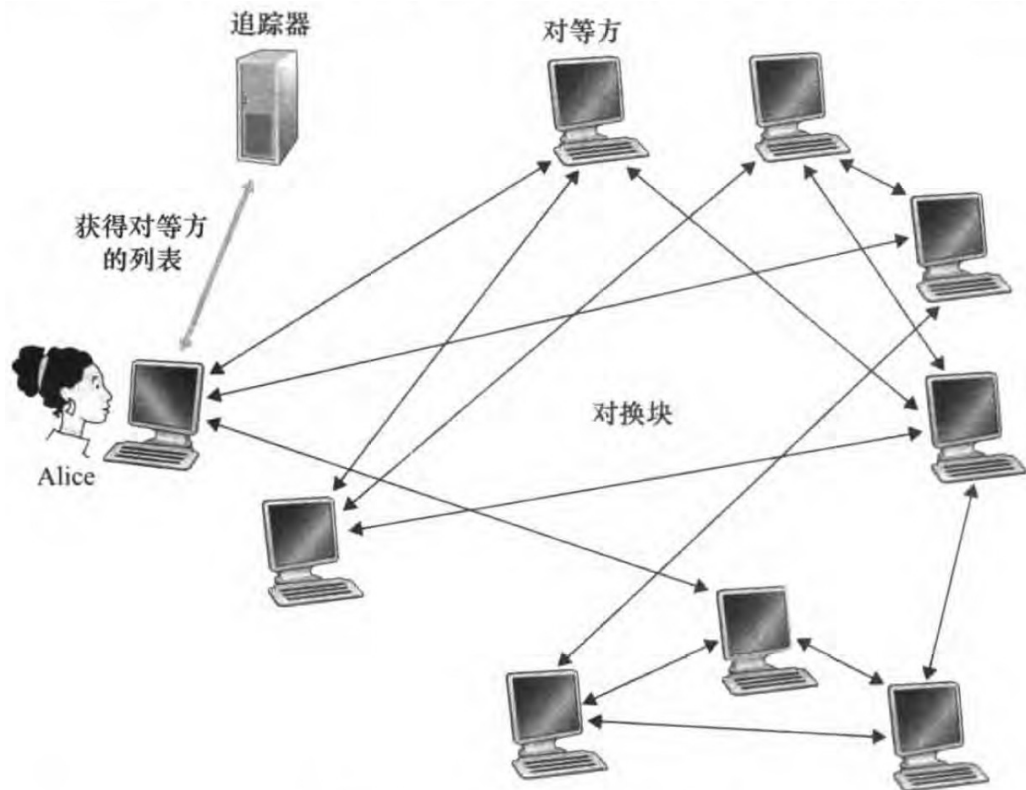


图 2-23 用 BitTorrent 分发文件





Internet Applications

- Internet Applications Overview
- Domain Name Service (DNS)
- Electronic Mail
- File Transfer Protocol (FTP)
- WWW and HTTP
- Content Distribution Networks (CDNs)





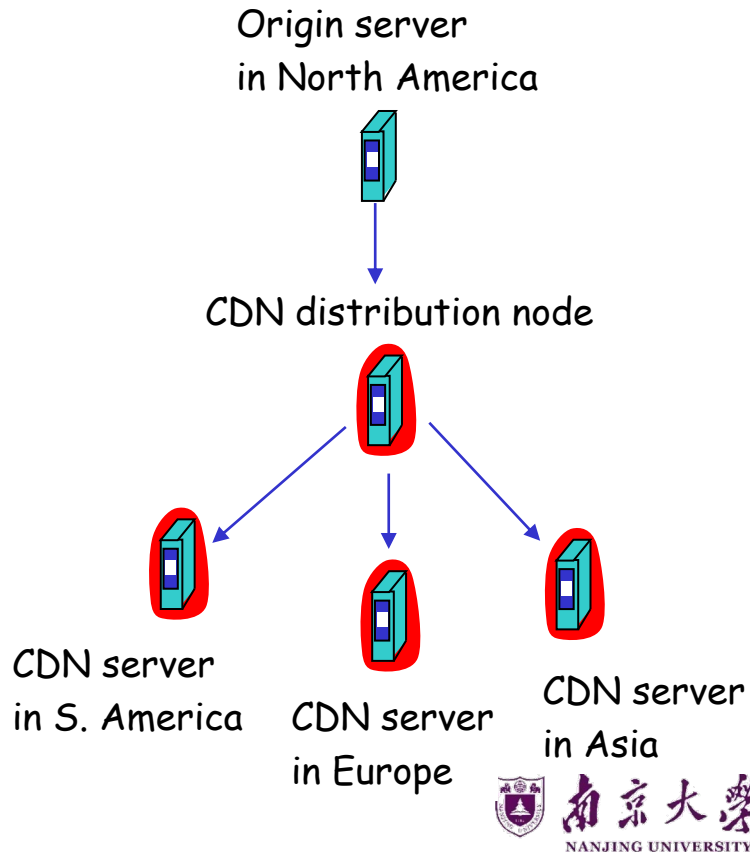
Content Distribution Networks (CDNs)

- **Challenge**

- Stream large files (e.g. video) from single origin server in real time
- Protect origin server from DDOS attacks

- **Solution**

- Replicate content at **hundreds of servers** throughout Internet
- **CDN distribution node** coordinate the content distribution
- Placing content **close to user**





Content Replication

- Content provider (origin server) is **CDN customer**
- CDN **replicates customers' content in CDN servers**
- When provider updates content, CDN updates its servers
- Use **authoritative DNS server** to redirect requests






Supporting Techniques

- DNS
 - One name maps onto many addresses
- Routing
 - Content-based routing (to nearest CDN server)
- URL Rewriting
 - Replaces "http://www.sina.com/sports/tennis.mov" with "http://www.cdn.com/www.sina.com/sports/tennis.mov"
- Redirection strategy
 - Load balancing, network delay, cache/content locality





CDN Operation



1' **URL rewriting** - get authoritative server

1. Get near CDN server IP address
2. Warm up CDN cache
3. Retrieve pages/media from CDN Server



课程习题（作业）——截止日期：3月11日晚23:59

- **课本110-115页**：第R3、R5、R16、P9、P22题
- 提交方式：<https://selearning.nju.edu.cn/>（教学支持系统）

教学支持系统

- ▾ 2025 Spring
 - 本科生一年级
 - 本科生二年级
 - 本科生三年级
 - 本科生四年级
 - 研究生一年级
 - 智能软件与工程学院

互联网计算-智软院

教师: 殷亚凤

课后习题

- 第1章-计算机网络和因特网
- 第2章-应用层

第2章-应用层

课本110-115页：第R3、R5、R16、P9、P22题

- 命名：学号+姓名+第*章。
- 若提交遇到问题请及时发邮件或在下一次上课时反馈。



课程习题（作业）——截止日期：3月11日晚23:59

R3. 对两进程之间的通信会话而言，哪个进程是客户，哪个进程是服务器？

R5. 运行在一台主机上的一个进程，使用什么信息来标识运行在另一台主机上的进程？

R16. 假定 Alice 使用一个基于 Web 的电子邮件账户（例如 Hotmail 或 Gmail）向 Bob 发报文，而 Bob 使用 IMAP 从他的邮件服务器访问自己的邮件。讨论该报文是如何从 Alice 主机到 Bob 主机的。要列出在两台主机间移动该报文时所使用的各种应用层协议。





课程习题（作业）——截止日期：3月11日晚23:59

P9. 考虑图 2-12，其中有一个机构的网络和因特网相连。假定对象的平均长度为 850 000 比特，从这个机构网的浏览器到初始服务器的平均请求率是每秒 16 个请求。还假定从接入链路的因特网一侧的路由器转发一个 HTTP 请求开始，到接收到其响应的平均时间是 3 秒（参见 2.2.5 节）。将总的平均响应时间建模为平均接入时延（即从因特网路由器到机构路由器的时延）和平均因特网时延之和。对于平均接入时延，使用 $\Delta/(1 - \Delta\beta)$ ，式中 Δ 是跨越接入链路发送一个对象的平均时间， β 是对象对该接入链路的平均到达率。

a. 求出总的平均响应时间。

b. 现在假定在这个机构 LAN 中安装了一个缓存器。假定命中率为 0.4，求出总的响应时间。

P22. 考虑向 N 个对等方分发 $F = 20\text{Gb}$ 的一个文件。该服务器具有 $u_s = 30\text{Mbps}$ 的上载速率，每个对等方具有 $d_i = 2\text{Mbps}$ 的下载速率和上载速率 u 。对于 $N = 10, 100$ 和 1000 并且 $u = 300\text{kbps}, 700\text{kbps}$ 和 2Mbps ，对于 N 和 u 的每种组合绘制出确定最小分发时间的图表。需要分别针对客户 - 服务器分发和 P2P 分发两种情况制作。





助教信息

| | | | |
|------|----|------------|-------------------------------|
| 甘世维 | 博士 | 计算机学院 | sw@smail.nju.edu.cn |
| 刘晓 | 博士 | 智能软件与工程学院 | 602024720002@smail.nju.edu.cn |
| 郭博文 | 博士 | 智能软件与工程学院 | bowen@smail.nju.edu.cn |
| 德斯别尔 | 硕士 | 智能如软件与工程学院 | 522024720002@smail.nju.edu.cn |





提问

Q & A

殷亚凤

智能软件与工程学院

苏州校区南雍楼东区225

yafeng@nju.edu.cn , <https://yafengnju.github.io/>



南京大學
NANJING UNIVERSITY